

AI Data Centers and U.S. Grids: A 2026–2030 Regional Forecast

Chaitanya Chandran Melnatami

Vista Ridge HS
Cedar Park, TX, USA

chaitanyamelnatami@gmail.com

Abstract—The rapid expansion of AI-driven data centers is projected to dramatically increase U.S. electricity demand and challenge regional power grids. By 2030, data center energy could reach ~600 TWh (10–12% of U.S. power). This growth will not impact all regions equally: the Eastern Interconnection e.g. Northern Virginia’s massive “Data Center Alley” and Texas (ERCOT) face the greatest strains, while the Western Interconnection sees significant but comparatively smaller impacts. Key factors such as grid infrastructure age, transmission constraints, climate, state policy incentives, and local data center density create these regional differences. An empirical forecast model (2026–2030) is developed to project data center energy use and peak loads for each interconnection, based on these drivers. The model indicates that without intervention, surging AI-related loads could stress legacy grid components and exacerbate peak demand e.g. ERCOT’s summer peak might nearly double from ~85 GW in 2023 to ~145 GW by 2030 and complicate decarbonization efforts in some regions. Data centers already account for ~26% of Virginia’s power use and that share may double by 2030. Major challenges for each grid are identified, and mitigation strategies are proposed prioritizing renewable energy expansion, grid modernization, policy reforms, and demand-side measures. Prior work (2017–2025) on data center energy is reviewed, and it is shown how this analysis extends beyond existing literature by providing a region-specific, grid-focused perspective. Results indicate that with proactive planning e.g. new high-efficiency capacity, transmission upgrades, flexible operations, and clean energy integration, the U.S. grid can accommodate the AI-driven surge in data center load while maintaining reliability and climate commitments.

Index Terms— *Artificial Intelligence, Data centers, Load Forecasting, Power Grid, Power Usage Effectiveness.*

I. INTRODUCTION

The U.S. electric grid is confronting a major challenge from the rapid rise of energy-intensive AI data centers. While efficiency advances kept data center energy use relatively stable at ~1.8–2% of the nation’s electricity through the 2010s [1], this changed around 2019 as compute-heavy AI workloads drove demand higher. By 2023, U.S. data centers consumed about 147 TWh (3.7% of national electricity), ending nearly a decade of flat global data center usage [2][4].

This paper examines how AI-centric data center growth (2026–2030) will impact the Eastern, Western, and Texas (ERCOT) grid regions. The Eastern Interconnection faces grid strain in areas like Virginia due to dense hyperscale clusters [7]. The Western region must balance slower AI data center growth with aggressive renewable goals and limited water [6][11]. Texas’s ERCOT grid sees rapid expansion from data centers and crypto mining, risking resource adequacy as flexible loads could reach 60 GW by 2031 [9][10].

This work projects regional data center energy use (2026–2030), identifies grid challenges, and proposes mitigation strategies. [1][4]. Paper structure: Section II gives background on U.S. grids and data center energy; Section III reviews recent and projected regional trends; Section IV describes the empirical forecast model; Section V discusses challenges and solutions by region; Section VI reviews prior research and unique contributions; Section VII concludes with grid planning and policy recommendations. This work offers actionable insights for utilities, grid operators, and policymakers to manage AI-driven demand while supporting reliability and sustainability [2][4].

II. BACKGROUND

U.S. Power Grids and Regional Context

The United States is served by three primary synchronous power grids (interconnections), each having distinct characteristics:

Eastern Interconnection: Covering about 39 states in the East and Midwest, the Eastern Interconnection is the nation’s largest grid by load, operated by PJM, MISO, SERC and others. Much of its transmission and distribution infrastructure is outdated and was not built for today’s concentrated data center demand [7][8]. Tax incentives and low electricity rates have drawn dense clusters, especially in Northern Virginia’s “Data Center Alley,” which accounted for roughly 25–26% of Virginia’s power use in 2023 [5][7]. Other major hubs are found in Georgia, North Carolina, Ohio, and Illinois. Supplying these loads strains infrastructure, as Dominion Energy warns transmission into Loudoun County is near

capacity [7]. Data center demand often overlaps with peak periods due to the region's hot summers and cold winters. Ambitious clean energy targets, like Virginia's 100% carbon-free goal by 2045, risk being undermined if rising data center demand is met with fossil generation [5][11].

Western Interconnection: Covering 11 states from the Rockies to the Pacific, the Western Interconnection transmits power over vast distances from remote sources (Wyoming wind, Arizona solar, Pacific Northwest hydro) to coastal cities. The region leads in renewable integration, with California reaching 35% renewables in 2021 and retiring natural gas plants to achieve a 100% clean electricity target by 2045 [6]. Key data center hubs include Silicon Valley, the Pacific Northwest, and emerging Desert Southwest locations like Phoenix, Arizona. Challenges include aligning data center demand with variable renewable supply, ensuring reliability during periods of low solar output, and managing cooling water requirements in arid regions. Large Western data centers rely on evaporative cooling towers, consuming 100–300 million gallons annually for a single 30 MW facility, raising concerns in drought-prone states [6][11]. Wildfire risks and extreme heat waves prompt utilities to implement Public Safety Power Shutoffs, sometimes forcing data centers onto backup generators for extended periods. Growth must be managed in tandem with grid resilience and decarbonization policies.

Texas Interconnection (ERCOT): Covering most of Texas and operating independently, ERCOT runs an energy-only market and has historically maintained a reserve margin near 15% via scarcity pricing signals. Texas's business-friendly policies, ample land, and low electricity prices (\$0.05–\$0.07/kWh for industrial users) have attracted data centers and crypto miners, making Texas the second-largest data center state with about 450 facilities as of 2023. AI-oriented data center growth is booming, especially in the Austin–San Antonio corridor and Dallas–Fort Worth area. ERCOT's total load was ~360 TWh in 2022 and could exceed 600 TWh by 2030, with data centers comprising 10–15% of consumption, up from ~2–3% in 2020 [9][10]. The combination of explosive growth and limited power import capability poses a resource adequacy challenge, as demonstrated during the February 2021 winter storm when high outages led to widespread blackouts. Tens of GW of nearly continuous new load may double ERCOT's summer peak by 2031 if all proposed projects proceed [10]. ERCOT must build generation or manage demand response internally; however, Texas leads the nation in wind and solar capacity, offering an opportunity to power new data centers with clean energy if transmission and storage infrastructure are expanded [9][10].

Data Center Energy Use Patterns and Efficiency

Since 2019, surging AI and cloud workloads have outpaced efficiency, causing U.S. data center energy to climb from ~150 TWh in 2022 to a projected ~500–600 TWh by 2030 [2][4][14]. Modern AI data centers operate at much higher utilization (60–80%) than traditional ones (historically <30%), with PUE improving modestly from ~1.5 (2020) to ~1.3 (2030) [1][8].

Energy use remains geographically concentrated: about 15 states account for ~80% of U.S. data center electricity, with Virginia exceeding 10% by 2020 and Dominion Energy reporting ~2.6 GW load in 2022 (24% of peak demand), rising to ~6 GW by 2035 [5][7][12]. Other major states are Texas, California, Oregon, and Washington. Grid impacts are highly localized, making region-specific analysis essential.

Utilities and grid operators are adapting: PJM raised 2030 forecasts after Virginia's interconnection requests [7][8], ERCOT launched a "Large Flexible Load" category in 2023 for data center and crypto loads with proposed new market mechanisms [9], and Oregon's 2025 law requires large new loads to fund grid upgrades [8][13]. Hyperscalers like Google are shifting workloads to match renewables, aiming to lower emissions [6].

In short, U.S. data center energy demand—driven by AI—grows rapidly and unevenly, putting pressure on specific regional grids. The following sections will quantify this growth and address management strategies for each interconnection.

III. ENERGY GROWTH TRENDS AND ANALYSIS (2026–2030)

Several reports project that by 2030, data centers could account for ~10% of total U.S. power consumption, up from ~3–4% today [2][4]. Table I summarizes a medium-case projection for U.S. data center electricity usage in 2030, broken down by interconnection (with 2023 values for reference).

Table I – Projected Data Center Electricity Use by 2030 (Medium Scenario)

Region	2023 Energy Use	2030 Energy Use (projection)	Share of U.S. DC Load (2030)
Eastern US	~70 TWh [5]	~310 TWh	~52% (= half of U.S.)
Western US	~30 TWh [5]	~120–140 TWh	~15%
Texas (ERCOT)	~50 TWh [9]	~150 TWh	~25%
United States (Total)	~150 TWh [2][4]	~600 TWh	100% (~10–12% of total)

U.S. data center annual energy consumption is projected to quadruple from ~150 TWh in 2023 to ~600 TWh in 2030, reaching approximately 10–12% of national demand [2][4]. Eastern and ERCOT regions are expected to comprise ~75% of the 2030 data center load. If AI adoption is slower and efficiency improves, data centers may use ~300–350 TWh in 2030 (~6% of U.S. load) [1]. Accelerated AI demand or stagnant efficiency could push usage above 700 TWh by 2030 [4][11]. Despite uncertainty, the trend is clearly upward.

Regional Patterns: Growth remains concentrated in specific regions:

Eastern Interconnection: Anticipated to supply about half of U.S. data center energy by 2030 (~300 TWh)—a 4x jump from ~120 TWh in 2023. States such as Georgia, North Carolina, Ohio, and Illinois are adding hundreds of MW of new capacity. By 2030, Eastern data centers could increase by 10–15 GW of constant load, potentially comprising 20–30% of peak demand for some utilities. Dominion Virginia's peak could reach 35–40% data centers by 2030, up from ~24% in 2022 [5][7]. Substantial grid infrastructure upgrades will be necessary.

Western Interconnection: Data center consumption is projected at ~120–140 TWh/year by 2030, up from ~30 TWh in 2020 (~4–5x growth). Growth is slower in California due to costs and local moratoria [11], while Oregon, Arizona, and Utah are experiencing expansion, with some new sites exceeding 200 MW. Even with improved efficiency, Western data center load is likely to at least double by 2030, potentially making up ~10–15% of Western grid demand (from ~5% in 2020). Integration with a grid transitioning to renewables will be a key challenge.

Texas (ERCOT): The fastest growth is forecasted here, with data centers expected to use ~150 TWh/year by 2030, up from ~50 TWh in 2023 (~3x increase). ERCOT's forecast indicates "large flexible loads" reaching 18 GW by 2030 [9][10]; if operated continuously, this equals ~158 TWh/year. By 2030, these loads could represent ~25% of ERCOT's peak demand and energy use (up from ~5% in 2020) [9][10]. This shift creates a 24/7 base load, potentially reducing ERCOT's planning reserve margin to near zero in extreme cases [10]. Market adjustments, new generation, and load flexibility will be required.




Fig 1: US Power Interconnections & Data Center Hubs

All scenarios indicate a significant increase in data center energy use by 2030. Adaptation in generation, transmission, and operations will be required by grid operators. An empirical model quantifying electricity demand and peak load contributions will be presented in the following section, forming the basis for impact analysis in Section V.

IV. FORECAST MODEL FOR DATA CENTER ENERGY LOAD

A projection model links IT capacity, PUE, and utilization to energy consumption and peak load including installed IT capacity, facility efficiency, utilization, etc. For a given region r and year t , define:

- $I_r(t)$ = Installed IT capacity (MW): the total power draw of all data center IT equipment if running at full load.
- $PUE_r(t)$ = Power Usage Effectiveness: ratio of total facility power to IT power (dimensionless; lower is more efficient).
- $U_r(t)$ = Average IT utilization: fraction of IT capacity actually used on average (0–1).

- $U_{peak,r}(t)$ = Peak utilization: fraction of IT capacity in use at the time of the grid's peak demand.

Then the annual energy consumption $E_r(t)$ and peak load $L_{peak,r}(t)$ attributable to data centers in region r can be estimated by:

$$E_r(t) = I_r(t) \times PUE_r(t) \times U_r(t) \times 8760 \quad (1)$$

$$L_{peak,r}(t) = I_r(t) \times PUE_r(t) \times U_{peak,r}(t) \quad (2)$$

Assumptions: The model is calibrated to 2022 data and then projected through 2030 with region-specific assumptions:

Installed IT Capacity (I_r): Baseline 2022 industry values estimates (e.g. ~8 GW Eastern, ~4 GW Western, ~3 GW ERCOT). Growth is added based on known expansion plans and trend extrapolation. In a medium scenario, I_r is assumed to grow ~10–15%/year in Eastern, ~8–12%/year in Western, and ~20%/year in ERCOT through 2030 (reflecting numerous large projects in Texas). This yields roughly a doubling of Eastern capacity, ~2.5x in Western, and ~3x in ERCOT by 2030, in line with forecasts [2][10].

Power Usage Effectiveness (PUE_r): Gradual PUE improvement over time is assumed. Many new hyperscale data centers achieve PUE ~1.2–1.3, but the fleet average is projected to improve from ~1.5 in 2022 to ~1.3 by 2030 as older sites are upgraded and new efficient designs come online. Hotter climates (Texas and Arizona) are expected to have slightly higher PUE (~1.3–1.35) due to greater cooling needs, whereas cooler regions (Pacific Northwest) could average ~1.2. No major breakthroughs (like widespread liquid cooling or server photonics) are assumed by 2030, consistent with industry outlooks [1][6].

Average utilization (U_r): Historically, average server utilization was low (~20% for enterprise data centers). Cloud and AI workloads are assumed to drive much higher utilization. U_r is projected to rise over time and is considered region-agnostic (since it depends more on workload type). Baseline 2022 U_r ~0.5 (50%) for hyperscale operators [1], reaching ~0.6–0.7 by 2030 as AI workloads (which often run at 70–90% utilization around the clock) constitute a larger share. This aligns with reports that AI training clusters are kept near fully busy to maximize expensive hardware use [4].

Peak utilization ($U_{peak,r}$): Data center loads tend to be flat or only moderately weather-sensitive. U_{peak} is assumed to be ~0.9 (90% of servers active) in Eastern and Western regions, reflecting some diversity or demand response, and U_{peak} ~1.0 in ERCOT, reflecting that many Texas data centers are expected to run continuously and ERCOT's peak (summer late afternoon) does not significantly curtail their usage. If demand-response programs are effective, U_{peak} could be lower (as data centers temporarily shed loads at peak), which is discussed in mitigations.

Sensitivity Analysis: Leading hyperscalers achieved PUE 1.1–1.2 in 2025 (Google: 1.10, Meta: 1.15, Microsoft: 1.18), but

fleet average remains ~ 1.4 due to legacy facilities. This model assumes air cooling with incremental efficiency gains. Widespread adoption of liquid cooling or server photonics could reduce PUE to ~ 1.05 by 2030, lowering projections by 15-20%. However, liquid cooling faces infrastructure costs and retrofit challenges, limiting near-term deployment.

Table II – Parameter Sensitivity (2030 Projections)

Parameter (Base Case)	-10% Variation (Low)	+10% Variation (High)	Total Impact Range	% of 2030 Load Impact
PUE (1.3)	545 TWh	655 TWh	± 60 TWh	$\pm 10\%$
Utilization (65%)	570 TWh	630 TWh	± 30 TWh	$\pm 5\%$
IT Growth (12%/yr)	520 TWh	680 TWh	± 80 TWh	$\pm 13\%$

This model reproduces known 2022 values (e.g. ~ 150 TWh total U.S. data center energy [2]) and yields 2030 estimates consistent with Table I earlier. It forms the basis for quantifying certain impacts: e.g. how much new peak capacity is needed in ERCOT (from Eq. 2), or how much energy could be saved by efficiency (from Eq. 1). Using this model, the next section explores, for each interconnection, the top challenges posed by data center growth and how they can be addressed.

V. CHALLENGES & MITIGATION STRATEGIES

Based on the above trends and model results, each major grid region faces specific challenges from AI data center expansion. We focus on two key challenges per region and discuss potential mitigation strategies for each. Some challenges (e.g. infrastructure needs) are common across regions, but the details and priorities differ.

A. Eastern Interconnection (Eastern U.S.)

Challenge E1: Transmission & Distribution Constraints in Data Center Hubs. The main bottleneck in the Eastern grid is delivering enough power to concentrated data center hubs, especially in Northern Virginia. High-voltage transmission lines and local substations in Loudoun and Prince William counties are nearing capacity due to the surge of data center connections. In 2022, Dominion Energy reported that several 500 kV transmission circuits into Loudoun County faced overload risks, potentially delaying new data center connections until upgrades were finished [7]. Similar constraints exist around Atlanta, GA, and Columbus, OH, where numerous 100+ MW interconnection requests strain the network. Without upgrades, these bottlenecks could hinder data center expansion or cause local reliability issues such as voltage drops or thermal overloads during hot weather.

Mitigation: Accelerate grid upgrades and modernize delivery networks. Utilities and grid operators are fast-tracking transmission and distribution improvements for data center growth in the Eastern region—such as PJM’s \$1 billion package for new lines and substations in Northern Virginia [8], with expedited permitting by Virginia regulators [7][12]. Distribution enhancements include high-capacity substations, transformer and feeder upgrades, and dedicated 115 kV feeds

for large campuses (e.g., Georgia). Advanced technologies like dynamic line ratings and power flow controls optimize capacity until new lines are ready. Demand response programs also help temporarily by shifting loads or using on-site generation during peak hours [7]. Substantial, timely investment in grid infrastructure is essential to support AI-driven load growth.

Challenge E2: Aging Infrastructure & Reliability Risks. Much of the Eastern interconnection’s infrastructure is decades old, and large constant loads accelerate wear and reduce margins. Transformers, underground cables, and switchgear in urban and suburban tech hubs are handling higher continuous currents than originally designed for, leading to overheating or overloading during peaks. For instance, a 40-year-old 138 kV cable serving a data center cluster may now operate near its thermal limit year-round, increasing risk of equipment failures and outages that could impact both data centers and local communities. These concentrated loads may also alter power flow patterns, necessitating updates to protection settings and voltage control schemes. Without targeted upgrades, a single utility equipment failure could disrupt significant loads, including critical IT services.

Mitigation: Prioritize modernizing grid infrastructure in data center hubs by replacing or upgrading transformers, cables, breakers, and protection systems [12]. Dominion’s 2023–2027 plan commits significant funding to such reliability upgrades in Northern Virginia [12]. Deploy advanced monitoring and automation (e.g., fiber-optic sensors, SCADA) to manage grid conditions and voltage. Data centers can enhance reliability by increasing redundancy (multiple feeds, on-site backup, storage) and participating in fast frequency response programs to stabilize the grid [14]. Incentivizing grid services and updating T&D infrastructure are vital for maintaining reliability amid rapid load growth.

B. Western Interconnection (Western U.S.)

Challenge W1: Generation-Load Mismatch & Renewable Integration. Issue: The Western grid is adding substantial solar and wind while retiring fossil plants, but data centers consume power constantly and do not natively respond to supply variability. This mismatch can result in high data center demand during low renewable output (e.g., at night or during wind lulls), requiring more gas generation or imports and raising CO₂ emissions unless mitigated [6][11]. Conversely, abundant renewable generation may be curtailed if data center loads are not flexible. The California Energy Commission (2025) warned that data center growth could boost natural gas use during peak net-load hours if demand is not shifted or paired with storage, hindering clean energy targets [11]. The core issue is aligning flat data center consumption with variable renewable production.

Mitigation: Integrate data centers into the clean energy transition through storage and flexible operations. (1) Use utility-scale and on-site batteries to shift renewables to cover

data center loads during low generation periods—California targets >10 GW of batteries by 2030 for evening discharge [6]. Some data centers cycle backup batteries daily for grid support. (2) Employ demand response by scheduling non-essential workloads when renewable output is high; cloud providers like Google already shift computing to cleaner times, and even 10–20% load flexibility flattens peaks. CAISO incentivizes participation through flexible demand programs [6]. (3) Pair data center growth with direct renewable procurement and smart siting, such as co-location near renewable-rich regions and on-site generation with storage for grid reliability [5]. Policies like time-of-use rates, demand response contracts, and clean energy standards for large users are key to aligning data center demand with renewable integration goals.

Challenge W2: Long-Distance Transmission & Siting Challenges. Issue: New or upgraded transmission is often required for Western data centers to access remote generation, but projects face complex permitting, environmental reviews, and local opposition. For example, the SunZia project took over a decade to permit due to route issues [15]. Local siting challenges also arise, such as concerns over noise, water use, and infrastructure strain in communities like Silicon Valley, central Washington, and Phoenix [11]. The main challenges are to ensure timely construction of bulk transmission for data center hubs and to minimize local impacts through careful site selection and design.

Mitigation: Streamline transmission development and adopt strategic data center siting to minimize conflicts. (1) Accelerate projects with federal and state policy support, leveraging “National Interest Electric Transmission Corridors” for faster approvals [15]. (2) Use existing corridors and advanced technologies like HVDC and high-capacity conductors to boost transmission with less land impact, as seen with underground HVDC along California’s I-5. (3) Site data centers in industrial areas, near substations, or on former plant sites to reduce community disruption; employ advanced cooling, including water-free or recycled water systems, to address noise and water issues [6]. Engage communities and offer local grid upgrades and economic benefits to secure support [5]. Overall, responsible grid expansion and smart siting will enable Western Interconnection to balance data center growth with environmental and social objectives.

C. Texas (ERCOT)

Challenge T1: Rapid Demand Growth vs. Supply Expansion. ERCOT faces unprecedented demand growth—about 5–6% per year through 2030, far exceeding its historical ~2% rate [10]. The main challenge is ensuring generation and grid capacity keep pace. ERCOT’s energy-only market relies on price signals to drive new capacity, but surging data center loads may cause signals to lag or spike excessively. In 2022–2023, record summer peaks and conservation appeals highlighted the risk; by 2030, summer peaks could regularly surpass the current firm generation capacity of ~90 GW

[9][10]. If resource additions lag, reserve margins may erode, increasing the likelihood of reliability events such as emergency load shedding during extreme weather.

Mitigation: Generation and storage additions are being accelerated through market reforms and state initiatives. (1) In 2023, the Public Utility Commission of Texas approved the Performance Credit Mechanism (PCM), incentivizing new dispatchable generation by rewarding availability during tight conditions [9]. PCM implementation by 2025, with sufficient revenue for 10–20 GW of new gas peakers or demand response by 2030, is crucial. (2) The Texas Energy Fund (S.B. 2627) will provide up to \$10 billion for backup generation and reliability projects [9]. Rapid deployment of this fund to build or upgrade peaking plants is critical, with ERCOT projecting a need for ~7–10 GW of new dispatchable capacity by 2030 [10]. (3) ERCOT’s renewable and storage pipeline includes over 20 GW of solar and 5 GW of batteries by 2025 [9]. Integration of these resources, supported by transmission expansions, is planned to meet new data center demand. Battery storage, set to exceed 10 GW by 2030, will help manage evening and winter peaks. (4) Demand response programs already enroll ~1.8 GW of “controllable load” (mainly crypto mines) [10]. Expansion to data centers, with backup generators and interruptible rate contracts, is encouraged. Overall, Texas is reforming its market, funding reliable generation, and deploying renewables and storage to maintain ERCOT’s reliability amid AI-driven demand.

Challenge T2: Minimal Import/Export Capability (Isolated Grid). ERCOT’s limited interconnections (<1 GW DC ties) mean it cannot rely on neighboring grids for support. As data center loads grow, this isolation heightens risk; ERCOT must manage large load swings and supply shortages alone. The 2021 winter crisis underscored the vulnerability, and rising constant loads could increase the risk of emergency load shedding unless internal reserves are strengthened.

Mitigation: ERCOT is enhancing reserves by considering new HVDC interconnections, expanded demand-side reserves, and backup generation. (1) New DC ties, such as a 1000 MW link to Louisiana or Oklahoma, are being discussed to provide emergency support; in 2022, a 2000 MW emergency-only tie to the Eastern grid was proposed [9]. (2) Texas plans to maintain a larger internal reserve margin, with the Energy Fund investing in grid-scale backup generators, aiming for a strategic reserve of several GW by 2030. (3) Emergency interruptible programs for data centers, allowing rapid curtailment, and leveraging onsite generation to feed power back to the grid during emergencies are being explored [14]. These measures strengthen ERCOT’s self-reliance and, combined with generation and transmission expansions, should enable the grid to handle growing AI loads while maintaining reliability standards.

VI. PRIOR WORK AND COMPARISON

From 2017–2025, studies showed data center energy use stayed flat with just a ~6% increase from 2010–2018 despite a

550% jump in compute instances [1]. This led to optimism that ongoing efficiency could contain energy demand. However, since 2019–2022, efficiency gains have slowed and AI-driven demand has surged, causing usage to rise. The IEA warned in 2023 that global data center energy could double by 2030 if AI growth continues rapidly, ending the previous plateau [4]. This analysis’s projections match IEA’s high scenario, indicating data center loads could be a major grid factor by 2030.

Recent national reports (McKinsey 2024, EPRI 2024, Lawrence Berkeley National Lab) forecast substantial increases in data center loads and emissions unless the grid mix improves [2][3][11]. McKinsey (2024) estimates U.S. data centers may reach ~570 TWh by 2030 (~12% of total demand), echoing this study’s medium scenario [2]. This analysis also highlights grid reliability risks, noting ERCOT’s reserve margin challenges and PJM’s transmission needs in Virginia topics rarely addressed in earlier literature that focused mainly on energy and carbon metrics. Prior literature highlights improving PUE and achieving “24/7 renewable energy” as key strategies [1][6]. Unlike previous national, global, or state-focused work [1][4][11], this study provides a region-specific comparison: Texas prioritizes capacity, the West targets renewable integration and water impacts, and the East addresses infrastructure expansion and cost allocation.

Economic Impact Analysis

Table III – Regional mitigation Strategies Comparison Projected

Region	Primary Challenge	Key Mitigation	Est. Cost	Timeline	Status
Eastern	Transmission capacity exhausted	Grid upgrades, Dynamic Line Rating	\$15-20B	2026-2030	In progress
Western	Renewable integration failures w/o Storage	Storage, demand response	\$8-12B	2026-2030	Proposed
ERCOT	Supply-demand balance (Reserve margin drops below 10% without new capacity)	Performance Credit Management, Generation, Storage	\$10-15B	2026-2029	Approved

Note: Costs include both public and private investment. Industrial users in data center zones may see 5-10% rate increases by 2030 without cost recovery mechanisms

These estimates assume proactive infrastructure investment. Delayed action could increase costs 30-50% due to emergency construction premiums and reliability events. In summary, this work confirms the rise in data center energy use and the need for efficiency and clean energy, while expanding the scope to grid infrastructure and tailored regional strategies—delivering a more holistic perspective than earlier research, which largely centered on IT efficiency or general energy/emissions analysis.

VII. CONCLUSION

The rapid expansion of AI-focused data centers is set to reshape U.S. power grids by 2030, potentially driving data centers to consume around 10% of national electricity—up from about 3–4% in 2020—with highest concentrations in the Eastern and Texas interconnections [2][4]. This growth will

directly affect generation needs, transmission infrastructure, and grid operations in key regions.

Empirical modeling indicates that Eastern Interconnection faces challenges in upgrading transmission and distribution systems to serve new load clusters, such as those in Virginia and Georgia, while meeting clean energy goals and ensuring equitable cost allocation. Utilities in the East and PJM have initiated grid upgrade programs; the continuation of these efforts and implementation of special tariffs for high-load users are essential. The Western Interconnection must integrate large, constant data center loads amid rising renewable energy penetration, requiring expanded energy storage, demand flexibility, and strategic siting to minimize water and local impacts. Ultimately, the intersection of digital and energy sectors via AI data centers presents both risks and opportunities. Success will depend on coordination among data center operators, utilities, and regulators. Proactive management can drive grid modernization and accelerate clean energy adoption, using data center demand to support renewable investment and justify transmission upgrades. Neglect could increase grid stress and emissions. Early, region-specific planning is crucial. Encouragingly, many stakeholders are pursuing the right strategies. By continuing to expand capacity, strengthen networks, improve efficiency, and leverage flexibility, reliable support for AI growth and clean energy progress can be achieved, turning potential strain into a catalyst for a smarter, greener grid.

REFERENCES

- [1] E. Masanet *et al.*, “Recalibrating global data center energy-use estimates,” *Science*, vol. 367, no. 6481, pp. 984–986, 2020.
- [2] B. Atamturk *et al.*, “The coming wave of data center demand: How AI and cloud are reshaping power needs,” *McKinsey Energy Insights*, Sep. 2024.
- [3] Electric Power Research Institute (EPRI), *U.S. Data Center Energy Consumption Projections*, Product ID: 3002031198, 2024.
- [4] International Energy Agency, “*Data Centres and Data Transmission Networks – Analysis*,” IEA, Oct. 2023.
- [5] Pew Research Center, “What we know about energy use at U.S. data centers amid the AI boom,” *Pew Research Center – Short Reads*, Oct. 24, 2025.
- [6] Chen, X., Wang, X., Colacelli, A., Lee, M., & Xie, L., “Electricity Demand and Grid Impacts of AI Data Centers: Challenges and Prospects,” *arXiv preprint arXiv:2509.07218*, 2025.
- [7] M. Turner, *Loudoun County, Virginia: Data Center Capital of the World — “A Strategy for a Changing Paradigm”*, Loudoun County Board of Supervisors, Oct. 20, 2025.
- [8] PJM Resource Adequacy Planning Department, *2025 Long-Term Load Forecast Report*, PJM Interconnection, Jan. 24, 2025.
- [9] ERCOT, “Report on Capacity, Demand and Reserves (CDR),” Dec. 2024.
- [10] Z. Skidmore, “Texas power demand could reach 218 GW by 2031, fueled by data centers,” *Data Center Dynamics*, Apr. 2025.
- [11] A. Cantú, “The insatiable energy demands of data centers could increase fossil fuel emissions in California,” *Capital & Main*, Oct. 1, 2025.
- [12] Dominion Energy, *2023 Integrated Resource Plan*, Richmond, VA, May 2023 (Appendix 5: Data Center Load).
- [13] Oregon HB 3546, 82nd Legislative Assembly – 2023, “Relating to energy load; establishing large energy consumer rate classes,” Jul. 2025.
- [14] Review of Challenges and Research Opportunities for Control of Transmission Grids,” *IEEE Xplore*, Document 10589406.